
Table of Contents

Introduction	1
BEGIN CODE	2
Load in the Cropped face database	2
Calculate the "average" face from all photos	3
Computing the SVD	3
Describing the Compact SVD	4
project each image onto basis	5
Explore the rank of r (see how many modes we really need)	6
Creating and exploring faces	8
The PCA process	9
Load in the uncropped face database	18
Calculate the "average" face from all photos	19
Computing the SVD	20
Explore the rank of r (see how many modes we really need)	20
Creating and exploring faces	21
The PCA process	21
Discussing PCA with the uncropped faces	26

Introduction

```
%{  
Written by Ian Good on 10/20/2020. For questions, please reach me at:  
iangood@uw.edu
```

The goal of this file is to explore singular value decomposition (SVD) and principal component analysis (PCA).

SVD and PCA are incredibly powerful tools for a data scientist and a good

understanding of their uses and limitations will pay dividends over one's career.

SVD is often used in data reduction where the original data is transformed in a way

to need significantly fewer dimensions to characterize. This is incredibly helpful in machine learning

where the dimensionality of data is often enormous.

PCA is often used to identify the variance and covariance of a dataset while also serving as a data reduction method.

By understanding the structure of the data you are working with, you may be able to extract incredibly valuable features, allowing for interpretations which may have been missed.

Note: This work covers and uses the compact SVD.

Two datasets were sourced for this project, both from the Yale Faces Project. A link to both can be found below:

<https://drive.google.com/drive/folders/1SQ77P5t5RUWCSucmk4jPFbufFMX8VrJG>

```
%}
```

BEGIN CODE

```
%Clear Previous Workspace and Figures
```

```
clear all  
close all  
clc
```

Load in the Cropped face database

```
tic  
% Specify the folder where the files live.  
myFolder = 'C:\Users\PhD\Documents\MATLAB\ME 564\faces\CroppedYale\';  
    %CroppedYale\yaleB31    or yalefaces\  
% Check to make sure that folder actually exists. Warn user if it  
    doesn't.  
if ~isfolder(myFolder)  
    errorMessage = sprintf('Error: The following folder does not  
    exist:\n%s\nPlease specify a new folder.', myFolder);  
    uiwait(wfaceSetdlg(errorMessage));  
    myFolder = uigetdir(); % Ask for a new one.  
    if myFolder == 0  
        % User clicked Cancel  
        return;  
    end  
end  
% Get a list of all files in the folder with the desired file name  
    pattern.  
filePattern = fullfile(myFolder, '**/*.pgm'); %search for whatever you  
    need. The **/ searches subfolders as well  
theFiles = dir(filePattern);  
N = length(theFiles);  
  
% Size of each picture in the cropped dataset  
m = 192;  
n = 168;  
  
avg = zeros(m*n,1); % start the average face array  
A = [];  
count = 0;  
  
for k = 1 : length(theFiles)  
    baseFileName = theFiles(k).name;  
    fullFileName = fullfile(theFiles(k).folder, baseFileName);  
    imageArray = imread(fullFileName);  
    %figure  
    %imshow(imageArray); % Display image.  
    %drawnow; % Force display to update immediately.
```

```

    if(size(imageArray,3)==1)
        M=double(imageArray);
    else
        M=double(rgb2gray(imageArray));
    end
    %pause(0.01);
    R = reshape(M,m*n,1);
    A = [A,R];
    avg = avg + R;
    count = count + 1;
end
avg = avg /count;

```

Calculate the "average" face from all photos

```

avgTS = uint8(reshape(avg,m,n));
figure(1), imshow(avgTS);
sgtitle('Average Face');

%{
The average face was taken from all photos in the uncropped dataset.
It
represents 2414 photos from the Yale Faces Dataset.
%}

```

Average Face



Computing the SVD

```

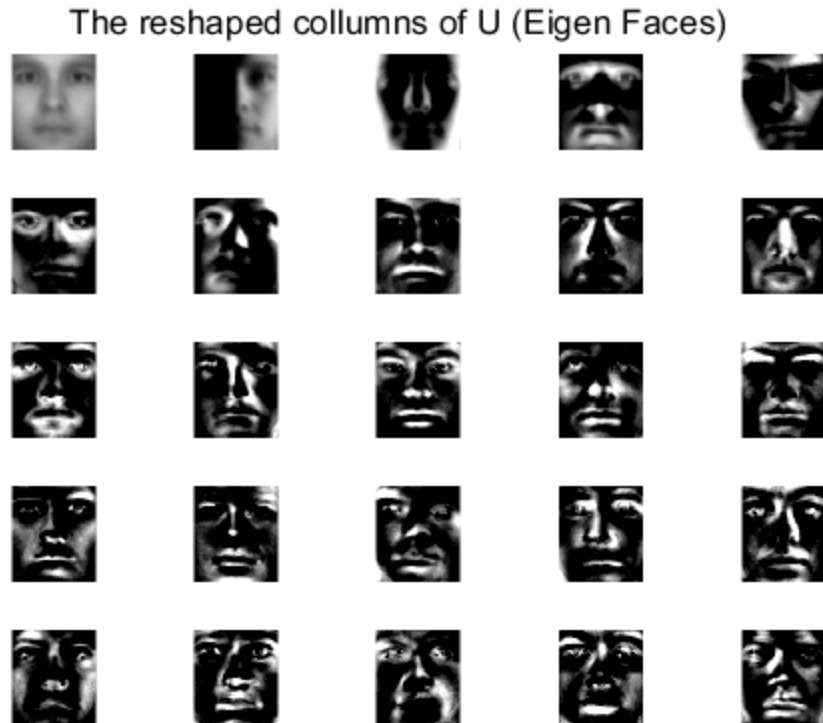
[U,S,V] = svd(A,0);
Phi = U(:,1:N);
Phi(:,1) = -1*Phi(:,1);
figure()
count = 1;
for iter=1:5

```

```

for j=1:5
    subplot(5,5,count)
    imshow(uint8(25000*reshape(Phi(:,count),m,n)));
    count = count + 1;
end
end
sgtitle('The reshaped collumns of U (Eigen Faces)');

```



Describing the Compact SVD

```

%{
The SVD consists of three output matrices U,S, and V*.
Given any complex matrix A (m x n) then...
U is an m x n unitary matrix which holds the left singular vectors of
A
S is an n x n diagonal matrix with the singular vales of A along the
diagonal arranged from highest to lowest with zeros everywhere else.
V is an n x n unitary matrix which contains the right singular vectors
of A

Since both U and V are unitary matrices, they form an orthonormal
basis and can be regarded as basis vectors.
For any linear transform T within the complex numbers of size n x m,
the svd maps as follows:
T(V_i) = sigma_i U_i, where sigma is the i-th diagonal entry of S.

```

The Range and Null space of a given matrix A can be represented using the SVD.

The right singular values (the V matrix) represents the null space and the left singular values span the Range.

From an intuitive perspective, the singular vectors from the SVD encode the semiaxes of an N-D ellipsoid while the singular values encode the magnitude of those semiaxes.

The Wikipedia page for SVD has a very beautiful transformation that displays this in 2D space:

<https://en.wikipedia.org/wiki/>

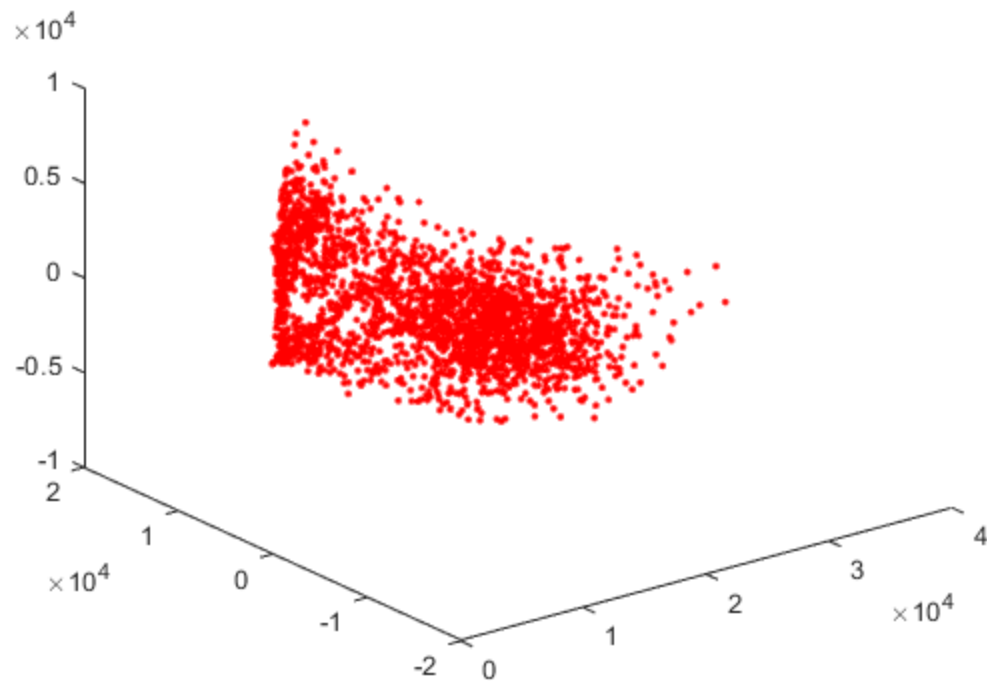
[Singular_value_decomposition#Geometric_meaning](https://en.wikipedia.org/wiki/Singular_value_decomposition#Geometric_meaning)

%}

project each image onto basis

```
for j = 1:N
    imvec = A(:,j);
    faceSet(:,j) = imvec'*Phi(:,1:3);
end
figure()
plot3(faceSet(1,:),faceSet(2,:),faceSet(3,:), 'r.')
hold on
title('Faces projected onto the first three left singular vectors');
```

Faces projected onto the first three left singular vectors



Explore the rank of r (see how many modes we really need)

```
fullModes = diag(S);
x = 1:length(fullModes);
figure();
subplot(1,2,1)
plot(x(1:32),fullModes(1:32)); %pick the first 32 modes so it is
    easier to visualise the highest importance modes.
xlabel('facial mode number');
ylabel('importance');
sgtitle('singular value spectrum (32 vs. full rank)');

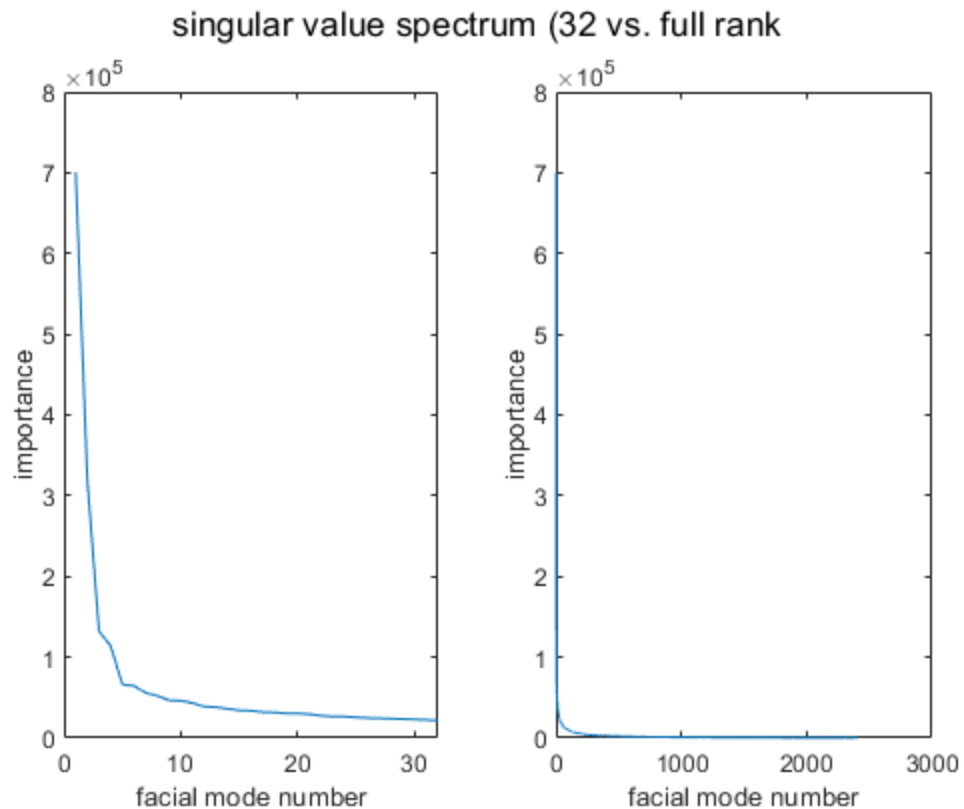
subplot(1,2,2)
plot(x,fullModes); %plot all of the modes so we can see the
xlabel('facial mode number');
ylabel('importance');

%{
As seen from the data, the modes take a long time to decay, indicating
the
rank of the dataset is large. We can also see the modes decay
following a
```

1/x relationship. This is often indicative of natural data which matches the dataset. While the first few values are larger, they are still within the same order of magnitude. If we had a truly dominant mode, we would expect it to be more than an order of magnitude above then next highest value.

It turns out that faces are rather complex structures, even when broken down into a low resolution photo. Within this datasett there are around two thousand, five hundred unique basis to evaluate a face. The fact that the human brain is able to so quickly and with little error is truly amazing!

The exact full rank of this dataset is: 2414
%}



Creating and exploring faces

```
if k == 2414 %number of photos in full cropped dataset. This is for
    compatability with the smaller set as there wouldn't be enough data
    to do this in the same way.

    hold on
    avgsIndFace = zeros(192,168,9); %Ind is individual
    avgAllFaces = zeros(192,168);
    figure()
    sgtitle('The average face of the first nine people in the
dataset')
    for people = 1:9 %pick nine people to compare (since that
fits nicely in a subplot). And the first nine datasets have 64
unncorrupted pictures each

        for j = 1:64 %photos per person

            imvec = A(:,j+(people-1)*64);
            avgsIndFace(:, :, people)
=reshape(imvec,m,n)+avgsIndFace(:, :, people);

        end
        avgsIndFace(:, :, people) =
uint8(1/64.*avgsIndFace(:, :, people));
        hold on

        subplot(3,3,people)
        imshow(avgsIndFace(:, :, people),[0,256]);

        avgAllFaces = 1/9.*avgsIndFace(:, :, people)+ avgAllFaces;

    end

    figure()
    imshow(avgAllFaces,[0,256])
    sgtitle('the average face of all nine photosets')
```

The average face of the first nine people in the dataset



the average face of all nine photosets



The PCA process

```
covariance = (A)*(A');  
[eigV,eigD] = eigs(covariance,100,'lm');  
  
for iter = 1:9  
    faceEigs = reshape(eigV(:,iter),m,n);
```

```

        subplot (3,3,iter)
        pcolor(flipud(faceEigs)), shading interp, colormap(gray),
set(gca,'Xtick',[],'Ytick',[])
    end
    sgtitle ('the first nine eigenvectors of the nine sets of
faces') %note how you can see the forming of the second eVal for the
full dataset (the half face).

    %now we reconstruct the faces from the eigenvalues and see how
close we get
    figure()
    sgtitle('The average single face projected onto the eigenvalues of
the dataset')
    for people = 1:9
        projFace = reshape(avgsIndFace(:, :, people), 1, 192*168)*eigV;
        subplot (3,3,people)
        bar(projFace(2:20)), set(gca,'Xlim',[0 20],'Ylim',[-2000
2000], 'Xtick',[], 'Ytick',[])
        text(0,-1450, strcat('Subject', num2str(people)), 'FontSize',
[15])

    end

    %now do the same for an individual face from each set. They should
have
    %similar eigenvalues to the average face from their set.

    figure()
    sgtitle('A single persons face projected onto the eigenvalues of
the average face')
    projSingFaceWeights = zeros(100,9);
    for people = 1:9

        projSingFaceWeights(:, people) =
reshape(A(:, 1+(people-1)*64), 1, 192*168)*eigV; %this picks the first
photo of each person, any can be chosen
        subplot (3,3,people)
        bar(projSingFaceWeights(2:20, people)), set(gca,'Xlim',[0
20], 'Ylim',[-2000 2000], 'Xtick',[], 'Ytick',[])
        text(0,-1450, strcat('Subject', num2str(people)), 'FontSize',
[15])

    end

    %plot the control faces
    figure
    sgtitle('Control faces for eigenface comparison')
    for people = 1:9
        subplot (3,3,people)
        imshow(reshape(A(:, 1+(people-1)*64), m, n), [0 256])
    end

```

```

    %plot the reconstructions of the faces from the eigenfaces

    eigenvaluesTested = [1,4,8,12,16,64,100];
    for iter = 1:length(eigenvaluesTested)

        eValsTested = eigenvaluesTested(iter);
        figure()
        sgtitle(strcat('The reconstructed faces using
',num2str(eValsTested), ' eigenvalue(s)'))

        for people = 1:9

            reconFace(:, :, people) =
eigV(:, 1:eValsTested)*projSingFaceWeights(1:eValsTested, people);
            subplot(3,3,people)
            imshow(reshape(uint8(reconFace(:, :, people)), 192, 168))

        end

    end

    %%PCA Discussion
    %{
        Plotting the face weights for the different subjects yielded
        results
        worth discussing. We see a significant difference between the
        average
        of a person's face versus their face in neutral lighting. It is
        suspected that this comes from the photos which have significant
        portions of the face shaded in darkness, bringing down the average
        pixel brightness, and requiring a significantly different
        combination of
        singular values to reconstruct. Subject 5, being the only one with
        bangs, had a very prominent identifier within the data. The sixth
        eigenface encoded bangs (as seen from the chart above) which
        allowed
        for easy identification.

        The benefits of looking at a range of subjects allows the
        extraction of
        data which might have been lost when looking just at a single
        person.
        This is a very prominent reminder that the algorithms we produce
        need
        to be trained with a wide variety of people, so as not to incur
        bias
        within the data.

        Looking to the reconstructed faces, we can get a feel for the rank
        of
        the face space. With just rank one, it is impossible to tell the
        subjects apart. With rank four, we start to see some changes in
        brightness but the subjects are largely unidentifiable. With rank
        8 we

```

start to see changes present (especially with subject five as we have included the 6th eigenface). With rank twelve, we can identify a few more faces but it is still difficult to see the difference between all subjects. At rank sixteen, all faces look distinct, but it would be difficult to exactly identify against the control images (especially for the leftmost column of subjects). At rank 64, all faces are clear and easily identifiable. There are still reconstruction errors present, especially around the eyes, but on the whole, the faces compare well against the control. At rank 100, there is even more detail present in the faces, with most incorrect blemishes removed. Errors still exist around the eyebrows and the camera flash is not present but this is of greater detail needed to identify the individuals.

Thus, we have taken the rank down from 2414 down to 64. This is a massive reduction in the required compute (37.7 times to be precise).

%}

Elapsed time is 818.173621 seconds.

else

%skip the face reconstruction since there isn't enough data here with a
%single person's dataset.

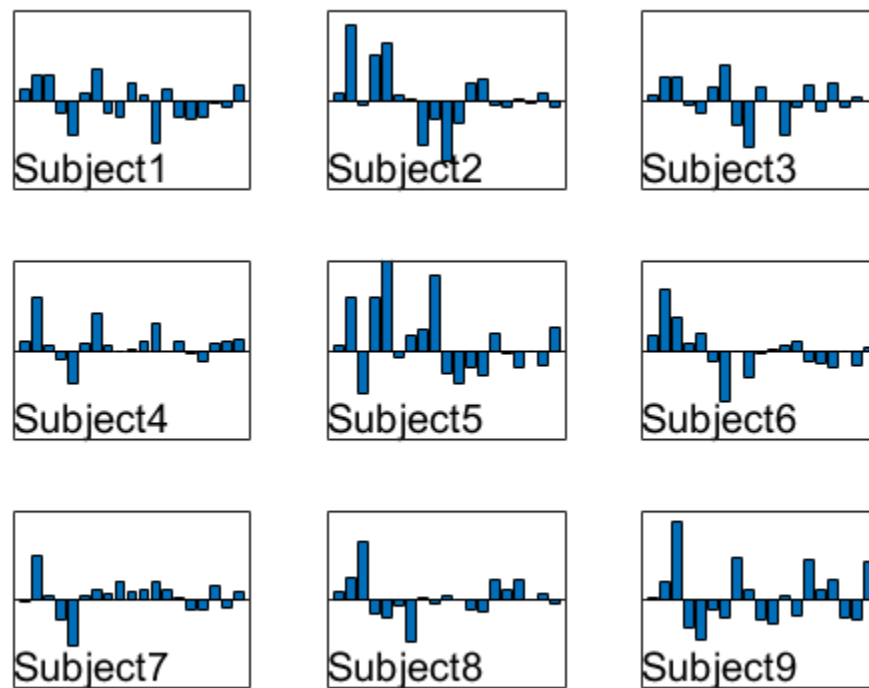
end

toc %to know how long the execute should take next time

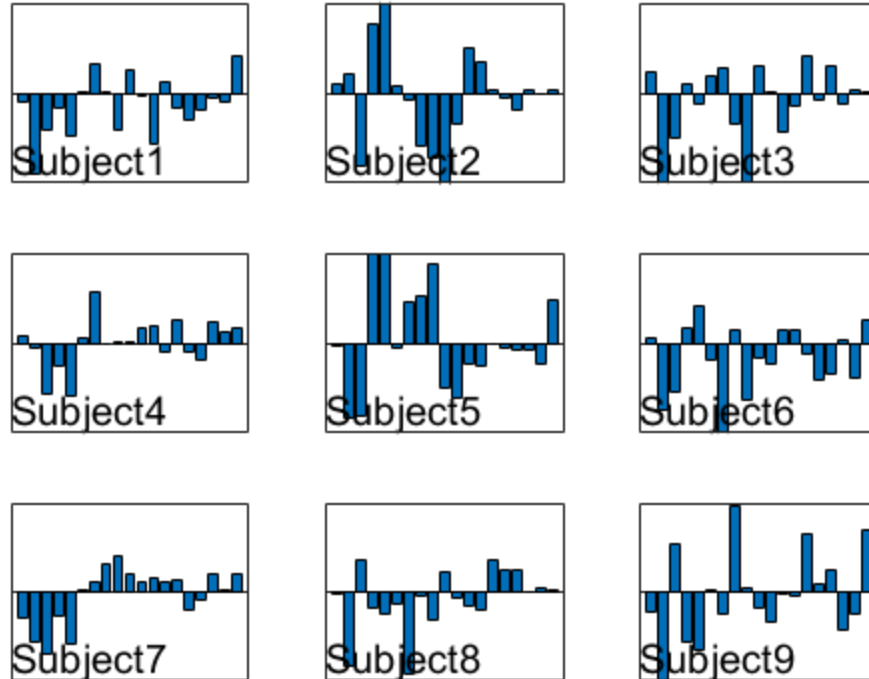
the first nine eigenvectors of the nine sets of faces



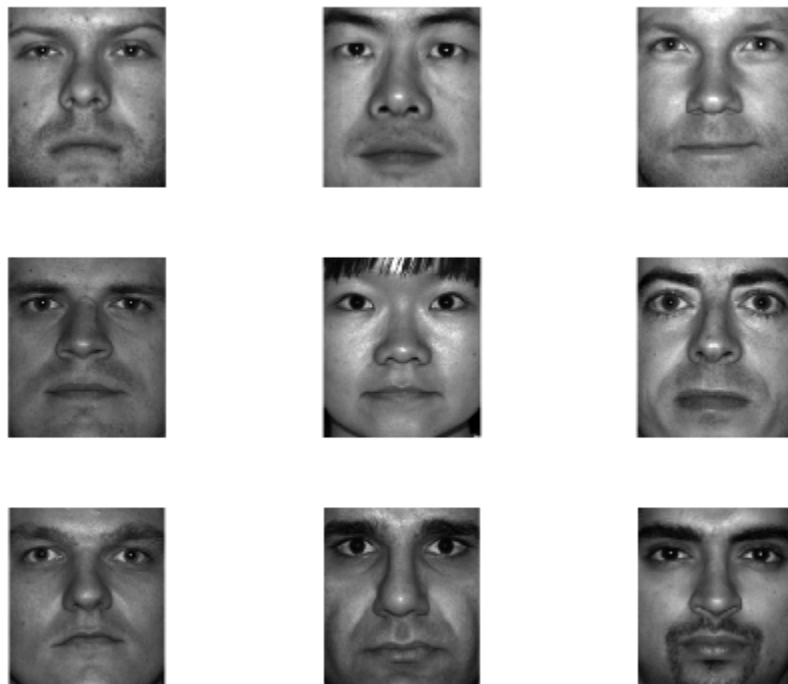
The average single face projected onto the eigenvalues of the dataset



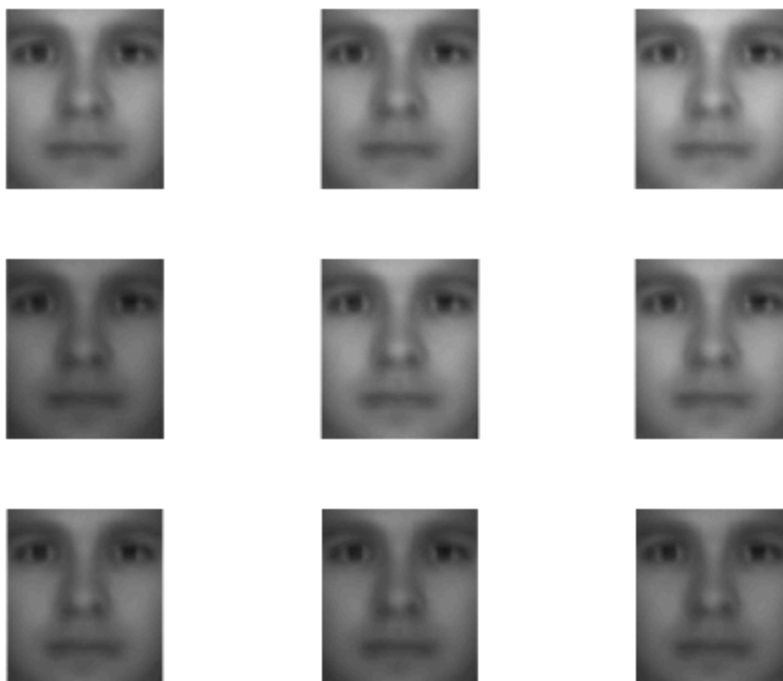
A single persons face projected onto the eigenvales of the average face



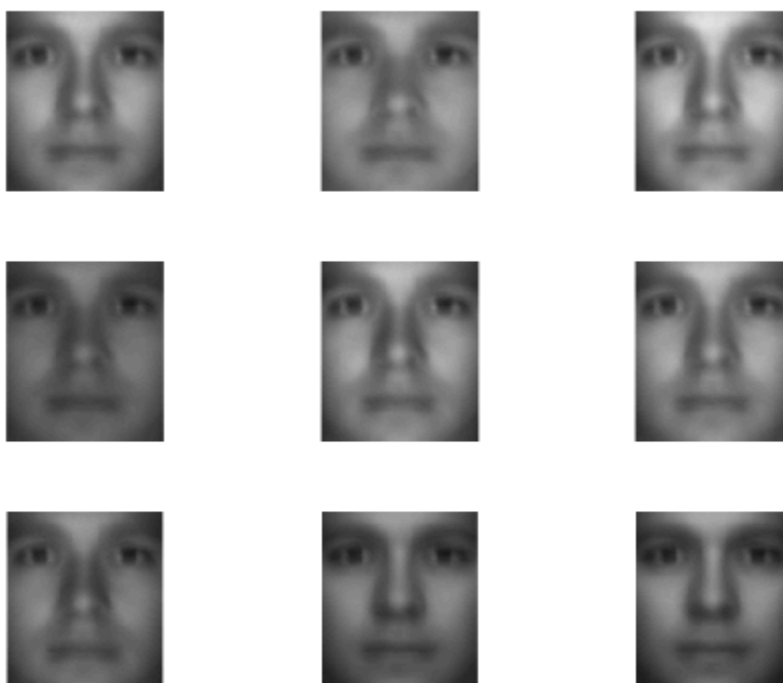
Control faces for eigenface comparison



The reconstructed faces using 1 eigenvalue(s)



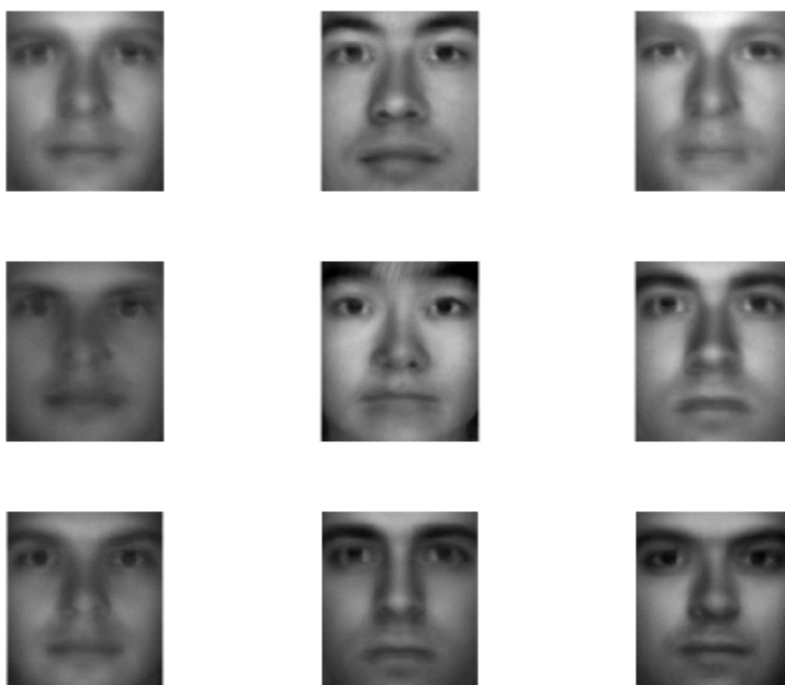
The reconstructed faces using 4 eigenvalue(s)



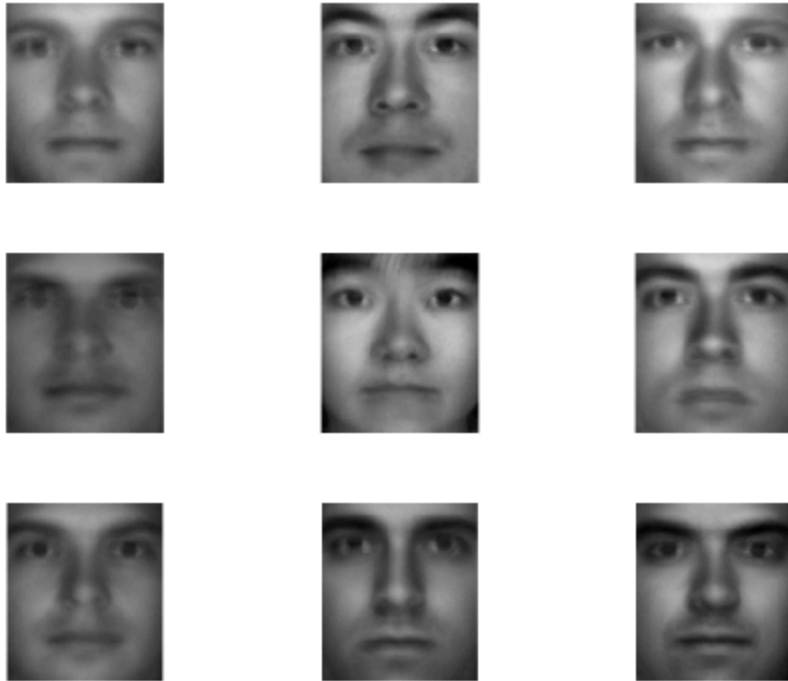
The reconstructed faces using 8 eigenvalue(s)



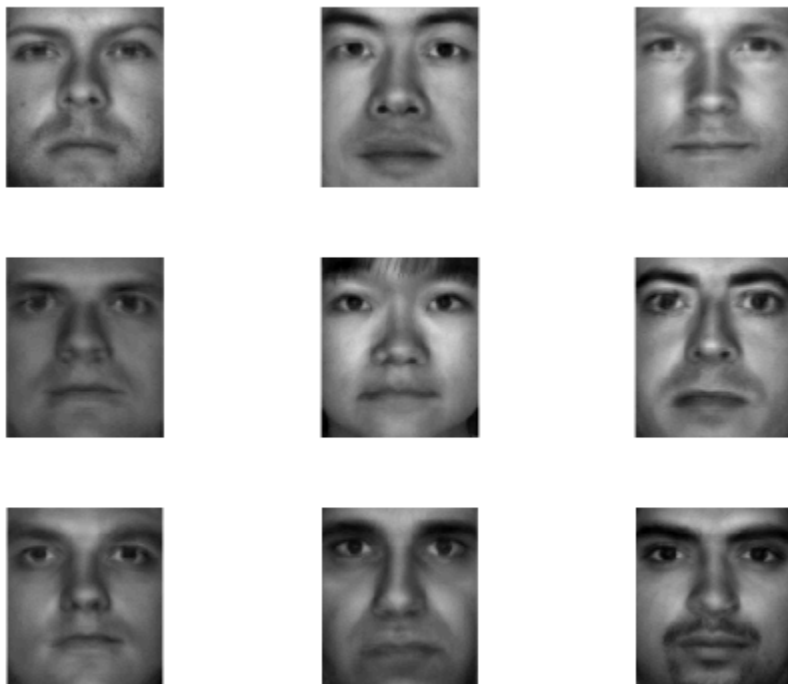
The reconstructed faces using 12 eigenvalue(s)



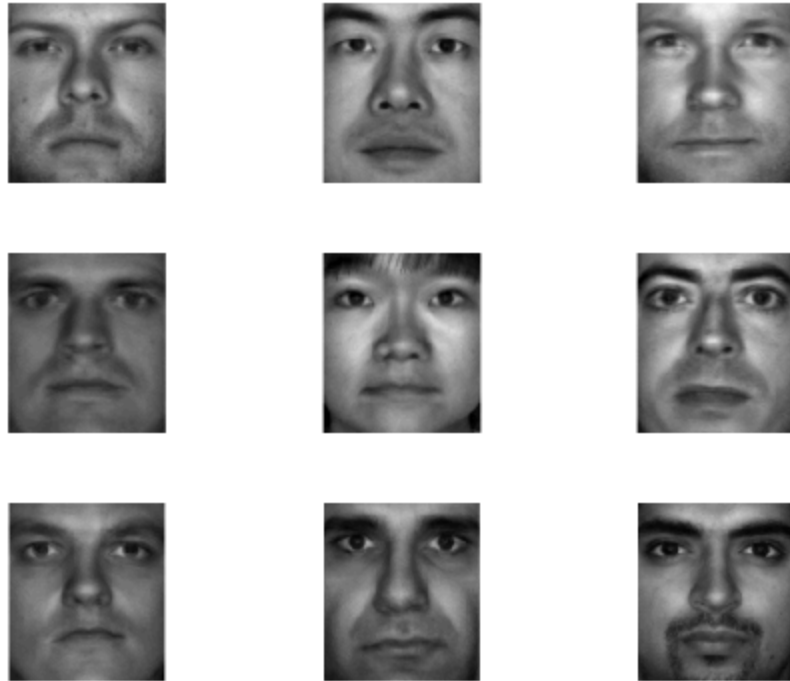
The reconstructed faces using 16 eigenvalue(s)



The reconstructed faces using 64 eigenvalue(s)



The reconstructed faces using 100 eigenvalue(s)



Load in the uncropped face database

```
% Specify the folder where the files live.
myFolder2 = 'C:\Users\PhD\Documents\MATLAB\ME 564\faces\yalefaces\';
%CroppedYale\yaleB31 or yalefaces\
% Check to make sure that folder actually exists. Warn user if it
doesn't.
if ~isfolder(myFolder2)
    errorMessage = sprintf('Error: The following folder does not
exist:\n%s\nPlease specify a new folder.', myFolder2);
    uiwait(wfaceSetdlg(errorMessage));
    myFolder2 = uigetdir(); % Ask for a new one.
    if myFolder2 == 0
        % User clicked Cancel
        return;
    end
end
% Get a list of all files in the folder with the desired file name
pattern.
filePattern2 = fullfile(myFolder2, '**/*.pgm'); %search for whatever
you need. The **/ searches subfolders as well
theFilesUncropped = dir(filePattern2);
N = length(theFilesUncropped);
```

```

% Size of each picture
m_UC = 180;
n_UC = 137;

avgUC = zeros(m_UC*n_UC,1); % the average face
A_UC = [];
count = 0;

for k = 1 : length(theFilesUncropped)
    baseFileNameUC = theFilesUncropped(k).name;
    fullFileNamUC = fullfile(theFilesUncropped(k).folder,
    baseFileNameUC);
    imageArrayUC = imresize(imread(fullFileNamUC),[n_UC
    m_UC]); %resize the imported photos to be easier to process in Matlab
%     figure
%     imshow(imageArrayUC); % Display image.
%     drawnow; % Force display to update immediately.

    if(size(imageArrayUC,3)==1)
        M=double(imageArrayUC);
    else
        M=double(rgb2gray(imageArrayUC));
    end
    %pause(0.01);
    R = reshape(M,n_UC*m_UC,1);
    A_UC = [A_UC,R];
    avgUC = avgUC + R;
    count = count + 1;
end
avgUC = avgUC /count;

```

Calculate the "average" face from all photos

```

avgTS_UC = uint8(reshape(avgUC,n_UC,m_UC));
figure(), imshow(avgTS_UC);
sgtitle('Average Face of the uncropped dataset');

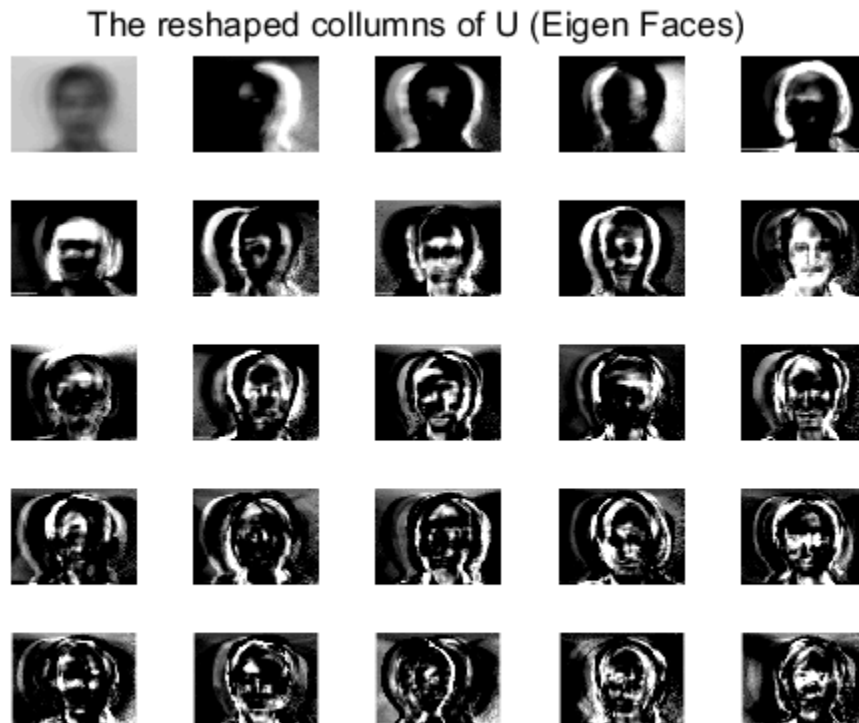
```

Average Face of the uncropped dataset



Computing the SVD

```
[U_UC,S_UC,V_UC] = svd(A_UC,0);
PhiUC = U_UC(:,1:N);
PhiUC(:,1) = -1*PhiUC(:,1);
figure()
count = 1;
for iter=1:5
    for j=1:5
        subplot(5,5,count)
        imshow(uint8(25000*reshape(PhiUC(:,count),n_UC,m_UC)));
        count = count + 1;
    end
end
sgtitle('The reshaped collumns of U (Eigen Faces)');
```



Explore the rank of r (see how many modes we really need)

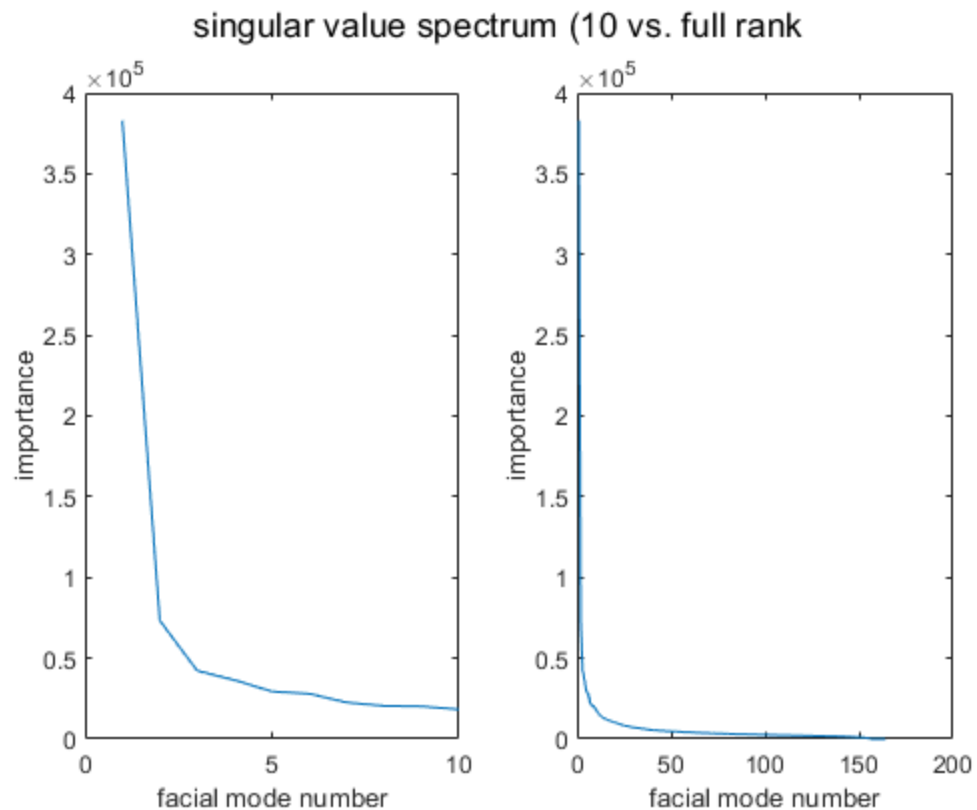
```
fullModesUC = diag(S_UC);
xUC = 1:length(fullModesUC);
figure();
subplot(1,2,1)
plot(xUC(1:10),fullModesUC(1:10));
```

```

xlabel('facial mode number');
ylabel('importance');
sgtitle('singular value spectrum (10 vs. full rank)');

subplot(1,2,2)
plot(xUC,fullModesUC);
xlabel('facial mode number');
ylabel('importance');

```



Creating and exploring faces

The PCA process

```

covarianceUC = (A_UC)*(A_UC');
[eigV_UC,eigD_UC] = eigs(covarianceUC,100,'lm');

```

```

%now do the same for an individual face from each set. They should
have
%similar eigenvalues to the average face from their set.

```

```

figure()

```

```

    sgtitle('A single persons face projected onto the eigenvales of
the average face') %but now with the weird ass faces
    projSingFaceWeights_UC = zeros(100,9);
    for people = 1:9

        projSingFaceWeights_UC(:,people) =
reshape(A_UC(:,5+(people-1)*11),1,n_UC*m_UC)*eigV_UC; %this picks the
first photo of each person, any can be chosen
        subplot (3,3,people)
        bar(projSingFaceWeights_UC(2:20,people)), set(gca,'Xlim',[0
20],'Ylim',[-2000 2000],'Xtick',[],'Ytick',[])
        text(0,-1450, strcat('Subject',num2str(people)), 'FontSize',
[15])

    end

    %plot the reconstructions of the faces from the eigenfaces

    eigenvaluesTested_UC = [1,4,16,32,64,100];
    for iter = 1:length(eigenvaluesTested_UC)

        eValsTestedUC = eigenvaluesTested_UC(iter);
        figure()
        sgtitle(strcat('The reconstructed faces using
',num2str(eValsTestedUC), ' eigenvalue(s)'))

        for people = 1:9

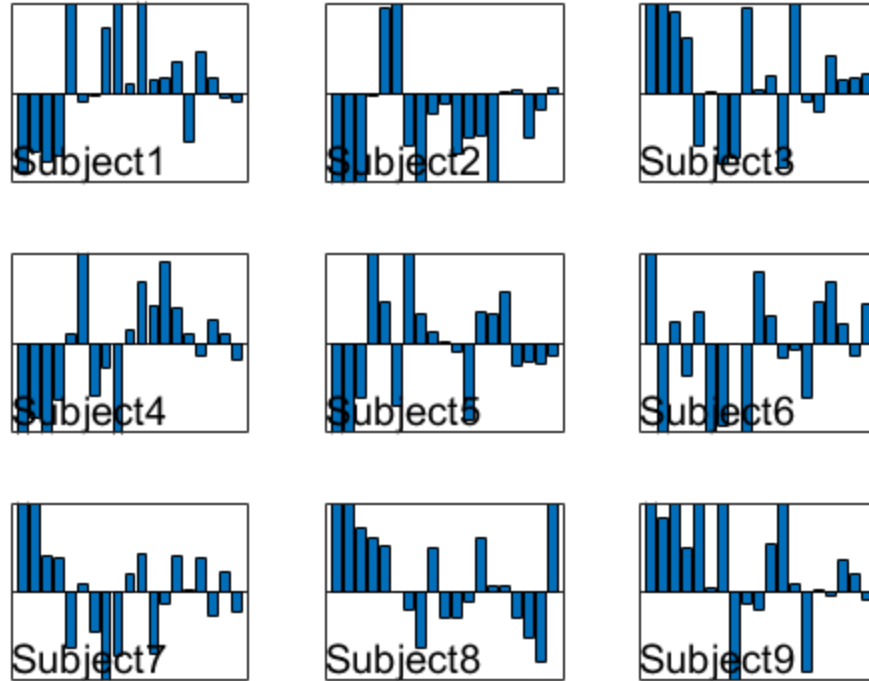
            reconFace_UC(:, :, people) =
eigV_UC(:,1:eValsTestedUC)*projSingFaceWeights_UC(1:eValsTestedUC,people);
            subplot(3,3,people)
            imshow(reshape(uint8(reconFace_UC(:, :, people)),n_UC,m_UC))

        end

    end
end

```

A single persons face projected onto the eigenvales of the average face



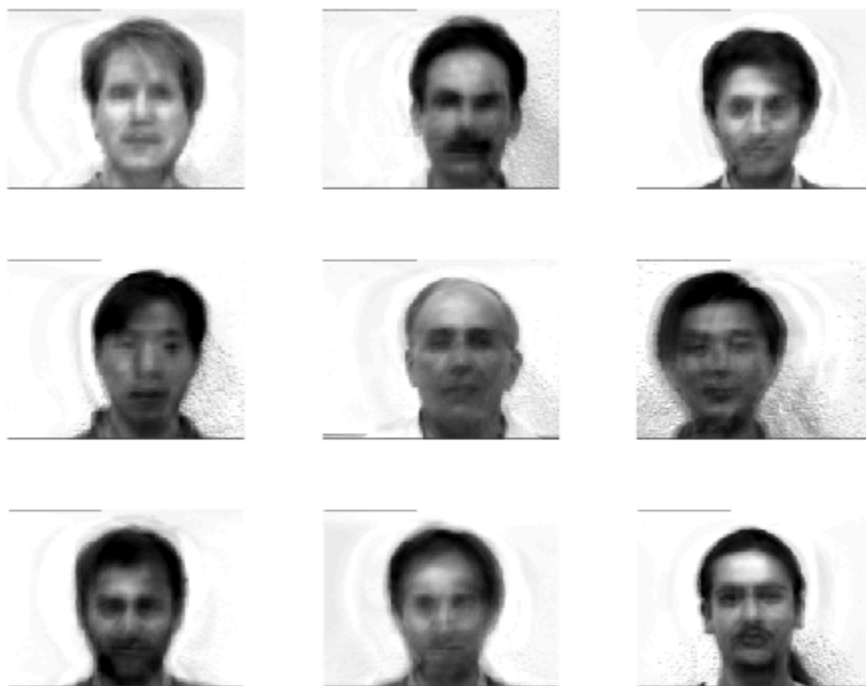
The reconstructed faces using1 eigenvalue(s)



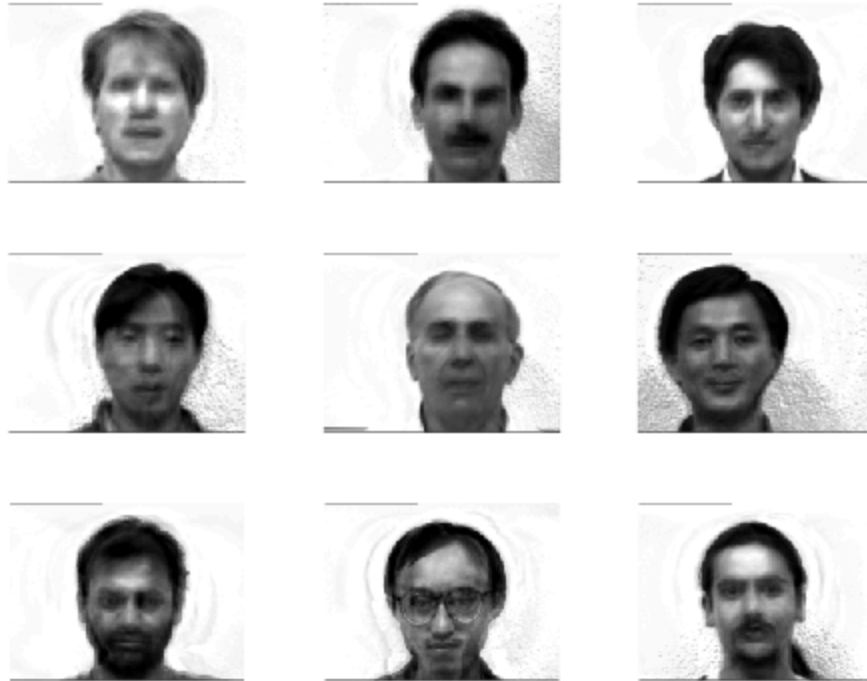
The reconstructed faces using 4 eigenvalue(s)



The reconstructed faces using 16 eigenvalue(s)



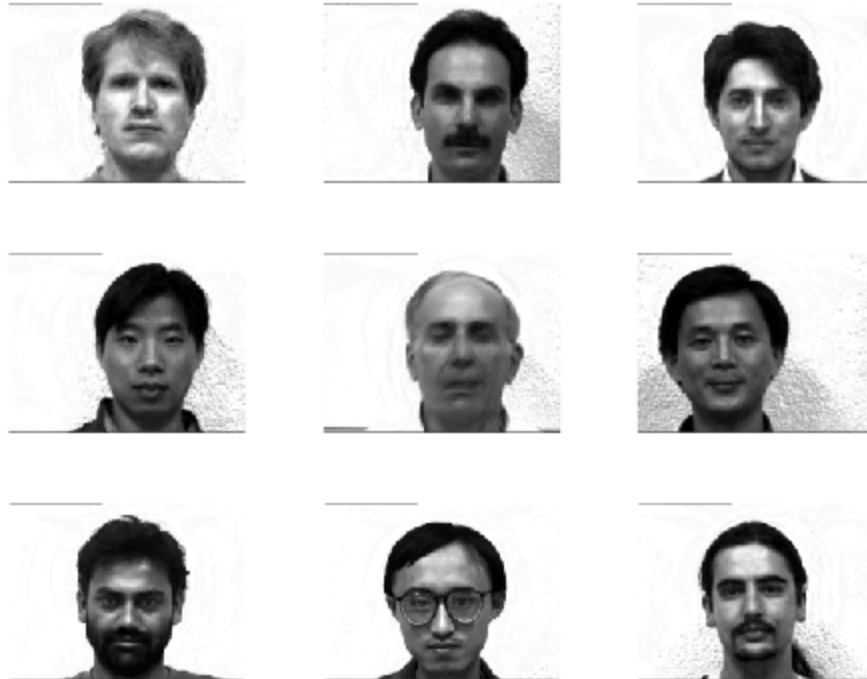
The reconstructed faces using 32 eigenvalue(s)



The reconstructed faces using 64 eigenvalue(s)



The reconstructed faces using 100 eigenvalue(s)



Discussing PCA with the uncropped faces

```
%{
    Looking to the reconstructed faces, we can get a feel for the
rank of
    the face space. With just rank one, it is impossible to tell the
subjects apart. With rank four, we start to see some changes in
brightness but the subjects are largely unidentifiable and full of
ghosting. Subject six is showing massive ghosting from the
different
positions their photos were taken from. Moving up to rank sixteen,
we
start to see discernable differences between subjects but most of
their
facial features are obscured or blurred. At rank 32, it is still
very
difficult to identify faces but features like hair color and hair
shape are discernable. There is noticeably less background ghosting
but
it is still very present within the reconstructions. At rank sixty
four
facial features are now identifiable but eyes and eyebrows are a
blurred mess. At rank 100, eyes and eyebrows are now separated for
all
subjects and would likely be identifiable. There is still a large
```

ammount of background ghosting present in the photos.

We also notice the decay of the singular values in this dataset was slower than that in the cropped and alligned data. We also see the singular values only stretch out to ~170. The fact that we needed most of the rank to have a halfway good chance at reconstructing the subject's faces meant we were not getting good data reduction. The set also follows a 1/x relationship telling us, it is likely made from natural sources.

Compared to the cropped and alligned images, this dataset required a much larger rank to recreate. Even at rank 100, the data was about equivalent to rank 16 with the cropped and alligned dataset. This shows the importance of properly setting up your data collection to get the desired results. There is a reason the saying garbage in, garbage out exists!

```
%}
```

```
%plot the control faces
figure
sgtitle('Control faces for eigenface comparison')
for people = 1:9
    subplot (3,3,people)
    imshow(reshape(A_UC(:,5+(people-1)*11),n_UC,m_UC),[0 256])
end
%%End Code
toc
```

Elapsed time is 950.256648 seconds.

Control faces for eigenface comparison



Published with MATLAB® R2019b